

Parallel and Distributed Computing -- (7A)		Name: <u>Nakeera Masood</u>
Quiz 03 (Fall 2022). Instructor: Dr. Syed M. Irteza		Roll Number: <u>19L-2201</u>
Date: 2022-10-11		
Total Marks: 15	Time Allowed: 10 mins	

- Assuming recursive doubling, what would the cost estimate of 1-to-all broadcast be (m = message size, t_s = startup time, t_w = transfer time, p = number of processes)? [2m]
 - $T = (t_s + mt_w) * \log(p)$
 - $T = (t_s + pt_w) * \log(m)$
 - $T = t_s + (pt_w * p)$
 - $T = (t_s + pt_w) * p$
- For a 2×16 mesh, using a naïve solution, how many messages would you expect the sending process to send, for a 1-to-all broadcast? [2m]
 - Less than 22
 - More than 30
 - Less than 10
 - Less than 16

$(2-1) \times 16$
 $2(16) \times$
- With all-to-1 reduction, we are essentially? [2m]
 - Requiring all $(p-1)$ processes to send messages to all p process
 - Requiring 1 process to send messages to $(p-1)$ processes
 - Requiring all $(p-1)$ processes to send messages to 1 process
 - Requiring all p processes to send messages to 1 process
- Can you describe the use of the following terms or clauses in OpenMP: [3*3m]
 - critical

critical section is the part of thread currently performing the code. It can be performing any set of instructions.

b. lastprivate(var)

The use of the ~~last~~ variable updated in the last process ~~is~~ when used in the next process is named lastprivate(var). It is basically using an updated variable in a thread.

c. schedule(static, 4)

static scheduling where the number of processes (threads used) is predefined. In this case it is 4.

4 is chunksize, not of threads.

⇒ critical clause has the critical section that has shared memory or variables. It is used to avoid the race condition and allow several variables to share a particular var of memory.

⇒ lastprivate (var)

assign last value of master thread to the variable given. This is a private var which is declared in OpenMP part and carries last value of master thread.

⇒ schedule (static, 4)

4 is chunk size.

Parallel and Distributed Computing -- (7A)	Name: <u>Moteeza Masood</u>
Quiz 04 (Fall 2022). Instructor: Dr. Syed M. Irteza	Roll Number: <u>19L-2201</u>
Date: 2022-10-25	
Total Marks: 15 (5*2 + 5)	Time Allowed: 10 mins

- Linked-List iteration was not something that could easily be parallelized in OpenMP due to:
 - We are unsure of the number of iterations involved
 - There is no parallel while loop construct in OpenMP
 - The data structure storing each node is very large
 - Both (a) and (b)
- When the chunk size is not specified, and we use static scheduling in OpenMP, the chunk size:
 - Equals 1
 - Equals total_threads
 - Equals $n/\text{num_of_procs}$ (where n = total iterations)
 - Equals $n/\text{total_threads}$ (where n = total iterations)
- With guided scheduling, the chunk size used:
 - Keeps increasing, till the size C (parameter) is reached
 - Keeps decreasing, till the size C (parameter) is reached
 - Stays constant
 - Decreases and increases within the range of size C (parameter)
- When OMP_DYNAMIC is set to true, then OpenMP will use:
 - The same number of threads as mentioned in the num_threads() clause
 - The same number of threads as mentioned in the omp_set_num_threads() function
 - The same number of threads as $[2 * \text{omp_get_num_procs}()]$
 - A dynamic adjustment algorithm to create threads that optimize system performance
- When we computed the value of π using the Monte Carlo method, we did not use the omp for construct. What feature of the problem enabled us to solve without needing that?

The #pragma omp parallel feature allowed us to compute the value of π .

- What could be the output of this program?

```
#include <iostream>
#include <omp.h>
using namespace std;
int main() {
    int nums[12] = { 13, 12, 11, 10, 9, 8, 7, 6, 5, 4, 3, 2 };

    #pragma omp parallel for num_threads(4) schedule(static, 3)
    for (int j = 0; j < 12; j++) {
        nums[j] += j;
        int x = omp_get_thread_num();
        cout << "At thread: " << x << " iteration: " << j << endl;
    }

    for (int i = 0; i < 12; i++) {
        cout << nums[i] << " ";
    }
    return 0;
}
```

R.W

num[j] = num[j] + j

num[j] = 13

j = 0

0 = 13 + 0

1 = 12 + 1

2 = 11 + 2

...

11 = 2 + 11

At Thread: 0 iteration: 1

0, 1, 2 3, 4, 5 6, 7, 8 9, 10, 11

Answers at back → (TPD)

Answers:

At	thread:	0	iteration:	0
At	thread:	0	iteration:	1
At	thread:	0	iteration:	2
"	"	1	"	3
"	"	1	"	4
"	"	1	"	5
"	"	2	"	6
"	"	2	"	7
"	"	2	"	8
"	"	3	"	9
"	"	3	"	10
"	"	3	"	11



13 13 13 13 13 13 13 13 13 13 13 13 13